

Bayesian Hierarchical Model for PM_{2.5} levels in Italy

Leila Cielok

29-04-2026

Contents

1	Introduction	2
2	Data	2
3	Model	4
3.1	Hierarchical structure	4
3.2	Prior specification	5
3.3	Posterior structure and conditional distributions	5
3.4	MCMC estimation	7
3.5	Implementation details	7
4	Results	8
4.1	Convergence diagnostics	8
4.2	Posterior inference	10
4.3	Posterior predictive checks	13
5	Conclusion	14

1 Introduction

Air pollution represents a major environmental and public health concern. Specifically, fine particulate matter $\text{PM}_{2.5}$ is among the most harmful pollutants due to its ability to penetrate deep into the respiratory system and its strong association with cardiovascular and pulmonary diseases. Reliable statistical modeling of $\text{PM}_{2.5}$ concentrations is therefore important for understanding pollution patterns and for assessing the persistence of environmental exposure across different locations.

A key challenge in modeling $\text{PM}_{2.5}$ concentrations lies in the presence of substantial heterogeneity across monitoring stations and irregular temporal coverage. Stations differ significantly in terms of geographical characteristics, local emission sources, and environmental conditions, while observational data often exhibit an unbalanced panel structure, with varying numbers of observations across stations and missing time periods. As a result, direct comparisons across stations may be misleading, since observed differences can reflect both genuine site-specific pollution patterns and differences in the amount and timing of available data.

In this context, Bayesian hierarchical models provide a natural and flexible framework to address these challenges. By introducing station-specific effects, they allow persistent differences across monitoring sites to be explicitly modeled, while still leveraging information from the entire dataset. This is particularly useful when some stations have limited observations, as it avoids relying exclusively on sparse local data and allows inference to be stabilized through partial pooling.

Building on these considerations, this paper develops a Bayesian hierarchical model for daily $\text{PM}_{2.5}$ concentrations observed across monitoring stations in Italy. Because the model includes a non-conjugate scale component, posterior inference is carried out using a Metropolis-within-Gibbs sampler, combining conjugate Gibbs updates with a Metropolis-Hastings step. The analysis focuses on model specification, MCMC implementation, convergence diagnostics, sensitivity analysis, and posterior predictive checks.

2 Data

The dataset comprises daily $\text{PM}_{2.5}$ concentration measurements collected across monitoring stations in Italy during 2025. The raw data were sourced from the European Environment Agency (EEA) Air Quality portal, specifically via the EEA Air Quality e-Reporting, with metadata and station specifications validated through the EEA Air Quality Viewer.¹ After preprocessing and quality control, the final dataset contains 31,349 observations distributed across 116 stations. Since the subsequent model is specified for a continuous, approximately Gaussian outcome, the distributional properties of $\text{PM}_{2.5}$ concentrations were inspected before estimation.

The empirical distribution of the raw outcome variable was highly right-skewed, with a skewness of approximately 17.1, primarily driven by extreme pollution events in the upper tail. Therefore, to stabilize variance and improve compatibility with the Gaussian likelihood of the proposed model, the outcome variable was transformed as

$$y_{it} = \log(\text{PM}_{2.5,it}),$$

where y_{it} represents the transformed concentration for station i at time t . This transformation

¹EEA Air Quality e-Reporting: <https://eadmz1-downloads-webapp.azurewebsites.net/>. EEA Air Quality Viewer: https://discomap.eea.europa.eu/App/AQViewer/index.html?fqn=Airquality_Dissemination.b2g.measurements.

substantially reduced skewness to approximately -0.29, yielding a more symmetric distribution, as shown in Figure 1, and making the likelihood specification a reasonable modeling choice.




Figure 1: Distribution of the log-transformed PM2.5 concentrations.

As is often the case in environmental monitoring networks, the resulting panel is markedly unbalanced. The number of observations per station ranges from 69 to 332, with a median of approximately 299. Temporal coverage is also uneven across stations. While most stations provide observations for most of the year, 32 stations lack at least one full month of data, and for a few stations coverage drops to only 4 months. This heterogeneity is summarized in Figure 2, which reports both the distribution of the number of observations per station and the distribution of the number of months observed per station.




Figure 2: Data availability by station.

For the estimation of the Bayesian hierarchical model, only a restricted set of variables was retained from the cleaned dataset. The modelling dataset includes the log-transformed outcome variable,

the station identifier, the month, and the day of the week. Observations with missing values in any of these variables were excluded from the modelling dataset, resulting in a complete-case sample. The variables `station_id`, `month`, and `day_of_week` were then treated as categorical factors. The corresponding design matrix includes an intercept, monthly indicators, and day-of-week indicators, while the station identifier is used to define the hierarchical grouping structure.

3 Model

The hierarchical nature of the resulting complete-case dataset, together with the strong imbalance in the panel structure, naturally motivates a multilevel modeling approach. In particular, the uneven number of observations across stations and the presence of incomplete temporal coverage make standard time-series or pooled regression models potentially inefficient and biased, as they are unable to properly account for unobserved heterogeneity across monitoring sites.

To address these issues, a Bayesian hierarchical specification is adopted, which allows for partial pooling across stations. This framework enables the model to borrow strength from the entire dataset while retaining station-specific components, thus providing robust inference even in the presence of sparse observations for some units.

The covariates included in the model correspond to calendar effects. In particular, the fixed-effects design matrix is constructed as

$$X = \text{model.matrix}(\sim \text{month} + \text{day_of_week}),$$

so that the intercept is included by default together with monthly and day-of-week indicators. The resulting parametrization corresponds to a treatment coding scheme, in which the intercept represents the expected log $\text{PM}_{2.5}$ level for the reference categories (January and Monday), and the remaining coefficients measure deviations from these baselines. Station-level heterogeneity is not included among the fixed effects, but is instead modeled through station-specific random intercepts.

The likelihood is specified as:

$$y_{it} \mid \beta, \alpha, \sigma^2 \sim \mathcal{N}(x_{it}^\top \beta + \alpha_i, \sigma^2),$$

where i indexes monitoring stations, t indexes the observation date, and x_{it} denotes the corresponding row of the fixed-effects design matrix X . The coefficient vector β captures the calendar fixed effects, while α_i is a station-specific random intercept accounting for time-invariant local characteristics.

Equivalently, the model can be written as

$$y_{it} = x_{it}^\top \beta + \alpha_i + \varepsilon_{it}, \quad \text{with } \varepsilon_{it} \sim \mathcal{N}(0, \sigma^2).$$

3.1 Hierarchical structure

Station-specific effects are modeled as draws from a common Gaussian population distribution:

$$\alpha_i \mid \tau \sim \mathcal{N}(0, \tau^2), \quad i = 1, \dots, J,$$

where J is the number of monitoring stations and τ^2 captures the cross-station variability in baseline pollution levels.

The assumption of a common Gaussian distribution reflects the idea that monitoring stations are exchangeable draws from a broader population of locations, once observed covariates are accounted for. This formulation allows the model to decompose variation in pollution levels into a systematic component (captured by $x_{it}^\top \beta$) and an unobserved station-specific component.

The hierarchical structure induces partial pooling: stations with limited observations are shrunk toward the overall mean, while stations with richer data retain more idiosyncratic estimates. This feature is particularly important in the present setting, where the panel is highly unbalanced and some stations provide only sparse temporal coverage. Shrinkage improves estimation stability and prevents overfitting to noisy local patterns.

3.2 Prior specification

Weakly informative priors are assigned to all parameters to regularize estimation while allowing the data to drive inference.

For the regression coefficients:

$$\beta \sim \mathcal{N}(b_0, B_0),$$

where $B_0 = cI_p$ is a diagonal covariance matrix. This specification assumes prior independence across coefficients and imposes mild shrinkage toward zero, helping to stabilize estimation in the presence of potentially correlated temporal indicators.

The observation variance is assigned the prior

$$\sigma^2 \sim \text{Inverse-Gamma}(a_\sigma, b_\sigma),$$

which provides a conjugate prior facilitating posterior computation while remaining weakly informative for appropriate choices of hyperparameters.

For the group-level standard deviation:

$$\tau \sim \text{Half-Cauchy}(0, A_\tau).$$

The use of a half-Cauchy prior for τ avoids the well-known sensitivity of inverse-gamma priors for variance components, while providing weakly informative regularization.

3.3 Posterior structure and conditional distributions

Combining the likelihood with the prior distributions, the joint posterior density is proportional to

$$p(\beta, \alpha, \sigma^2, \tau \mid y) \propto \left[\prod_{i,t} \mathcal{N}(y_{it} \mid x_{it}^\top \beta + \alpha_i, \sigma^2) \right] \left[\prod_{i=1}^J \mathcal{N}(\alpha_i \mid 0, \tau^2) \right] \mathcal{N}(\beta \mid b_0, B_0) \cdot \text{IG}(\sigma^2 \mid a_\sigma, b_\sigma) \cdot \pi(\tau).$$

Although the joint posterior does not admit a closed-form expression, most full conditional distributions are available in standard form, while the scale parameter τ does not admit a conjugate full conditional distribution under the half-Cauchy prior and therefore requires a non-standard update.

In particular, conditional on σ^2 and τ , the regression coefficients β follow a multivariate normal distribution:

$$\beta \mid \cdot \sim \mathcal{N}(\tilde{b}, \tilde{B}),$$

where the posterior mean and covariance matrix are given by

$$\tilde{B} = \left(B_0^{-1} + \frac{1}{\sigma^2} X^\top X \right)^{-1}, \quad \tilde{b} = \tilde{B} \left(B_0^{-1} b_0 + \frac{1}{\sigma^2} X^\top (y - \alpha) \right),$$

with y denoting the stacked outcome vector and α the vector of station-specific effects appropriately expanded to match the observation-level dimension.

Similarly, each station-specific effect admits a Gaussian full conditional:

$$\alpha_i \mid \cdot \sim \mathcal{N}(\tilde{a}_i, \tilde{V}_i),$$

where

$$\tilde{V}_i = \left(\frac{n_i}{\sigma^2} + \frac{1}{\tau^2} \right)^{-1}, \quad \tilde{a}_i = \tilde{V}_i \cdot \frac{1}{\sigma^2} \sum_{t \in \mathcal{T}_i} (y_{it} - x_{it}^\top \beta),$$

with n_i denoting the number of observations available for station i , and $\mathcal{T}_i$ the corresponding set of time indices.

The observation variance retains conjugacy and follows an inverse-gamma full conditional:

$$\sigma^2 \mid \cdot \sim \text{Inverse-Gamma}(\tilde{a}_\sigma, \tilde{b}_\sigma),$$

where

$$\tilde{a}_\sigma = a_\sigma + \frac{N}{2}, \quad \tilde{b}_\sigma = b_\sigma + \frac{1}{2} \sum_{i,t} (y_{it} - x_{it}^\top \beta - \alpha_i)^2,$$

and N denotes the total number of observations.

In contrast, the full conditional distribution of τ does not belong to a standard family due to the half-Cauchy prior. Its conditional density is proportional to

$$p(\tau \mid \alpha) \propto \left[\prod_{i=1}^J \mathcal{N}(\alpha_i \mid 0, \tau^2) \right] \pi(\tau), \quad \tau > 0,$$

and is therefore updated via a Metropolis-Hastings step within the Gibbs sampler.

These conditional distributions directly motivate the Metropolis-within-Gibbs algorithm described in the next section.

3.4 MCMC estimation

Posterior inference is performed via a Metropolis-within-Gibbs sampler. Conditional on the variance and scale components σ^2 and τ , the regression coefficients β and the vector of station-specific effects $\alpha = (\alpha_1, \dots, \alpha_J)$ admit Gaussian full conditional distributions and are therefore updated through Gibbs steps. The observation variance σ^2 is sampled from its inverse-gamma full conditional distribution.

The group-level scale parameter τ , for which no closed-form full conditional is available under the half-Cauchy prior, is updated through a random-walk Metropolis step. This hybrid scheme combines the computational convenience of conjugate updates with the flexibility required for the non-conjugate scale component.

The algorithm therefore alternates between Gibbs updates for β , α , and σ^2 , and a Metropolis-Hastings update for τ .

3.5 Implementation details

Hyperparameters are selected to induce weakly informative priors. For the regression coefficients, the prior covariance matrix is set to $B_0 = cI_p$, with baseline choice $c = 100$, implying a large prior variance and only mild shrinkage toward zero. For the observation variance, the inverse-gamma hyperparameters are fixed at $a_\sigma = 2$ and $b_\sigma = 0.5$, yielding a diffuse prior. The value $b_\sigma = 0.5$ is chosen to be consistent with the empirical scale of the log-transformed outcome (whose sample variance is approximately 0.47), while remaining sufficiently weakly informative. The scale parameter of the half-Cauchy prior on τ is set to $A_\tau = 5$, allowing for substantial between-station variability while still providing regularization. The robustness of posterior inference to alternative hyperparameter choices is examined through a dedicated sensitivity analysis reported below.

Posterior inference is performed using a Metropolis-within-Gibbs sampler run for 10,000 iterations. Burn-in and thinning are applied as post-processing steps. The burn-in period is selected using the Heidelberger–Welch diagnostic, which suggests discarding approximately 1,000 initial iterations based on stationarity considerations. Accordingly, the first 1,000 iterations are removed, and every 4th remaining draw is retained to reduce storage and autocorrelation.

The Metropolis step for τ is implemented as a Gaussian random walk proposal on $\log \tau$. After tuning, a proposal standard deviation of 0.20 is adopted, yielding an acceptance rate of approximately 0.37. Posterior summaries and diagnostic plots are essentially unaffected by alternative choices, so this specification is retained for the final MCMC run.

Finally, convergence is assessed using both graphical (trace plots and autocorrelation functions) and numerical diagnostics, including the Geweke statistic, the Heidelberger–Welch diagnostic, and the effective sample size.

3.5.1 Sensitivity to prior specification

To assess whether posterior inference is materially affected by the baseline prior calibration, a one-at-a-time sensitivity analysis is performed. Starting from the baseline configuration $c = 100$, $A_\tau = 5$, and $b_\sigma = 0.5$, each hyperparameter is varied separately while keeping the others fixed.

For computational efficiency, the sensitivity runs are based on shorter auxiliary chains of 5,000 iterations, with the first 200 discarded as burn-in and the same thinning rule applied as in the baseline analysis. Since the purpose of this exercise is to assess the stability of posterior summaries rather than to obtain the final inferential estimates, this reduced MCMC setting is considered adequate.

Three prior components are examined. First, the prior variance of the regression coefficients is changed by setting $c = 10$ and $c = 1000$. Second, the scale parameter of the half-Cauchy prior on the station-level standard deviation is varied from $A_\tau = 5$ to $A_\tau = 2.5$ and $A_\tau = 10$. Third, the inverse-gamma scale parameter for the observation variance is increased from $b_\sigma = 0.5$ to $b_\sigma = 1$, while keeping $a_\sigma = 2$ fixed.

As reported in Table 1, posterior summaries are stable across all specifications. The posterior means of the intercept, σ^2 , and τ remain virtually unchanged, and the substantive interpretation of the temporal effects and station-level heterogeneity is unaffected. Acceptance rates for the Metropolis step on τ also remain in a satisfactory range across specifications.

Table 1: Sensitivity analysis under alternative prior specifications.

Scenario	c	A_τ	b_σ	β_0	σ^2	τ	Acceptance (τ)
Baseline	100	5.0	0.5	2.603	0.325	0.299	0.388
Smaller prior variance for β	10	5.0	0.5	2.602	0.325	0.298	0.394
Larger prior variance for β	1000	5.0	0.5	2.603	0.325	0.299	0.401
Smaller A_τ	100	2.5	0.5	2.603	0.325	0.299	0.384
Larger A_τ	100	10.0	0.5	2.603	0.325	0.299	0.388
Larger b_σ	100	5.0	1.0	2.603	0.325	0.298	0.393

Overall, the sensitivity analysis indicates that posterior inference is not materially driven by the specific prior calibration. Instead, the likelihood appears sufficiently informative relative to the prior.

4 Results

4.1 Convergence diagnostics

Trace plots of the post-burn-in chains for the variance and scale components (Figure 3) show no evidence of non-stationarity or systematic trends. Both chains fluctuate around stable levels, indicating satisfactory convergence of the Metropolis-within-Gibbs sampler.




Figure 3: Post-burn-in trace plots for the variance and scale components.

Autocorrelation functions for the variance and scale components are reported in Figure 4. For both σ^2 and τ , autocorrelation declines rapidly after the first lag and remains close to zero thereafter, indicating efficient mixing for the variance components.




Figure 4: Post-burn-in autocorrelation functions for the variance and scale components.

Autocorrelation functions were also inspected for representative regression coefficients and station-specific effects, selected according to their effective sample size in order to cover parameters with low, median, and high sampling efficiency. While some persistence is observed at the first lags, especially for a subset of station-specific effects, autocorrelation decreases progressively and does not display long-range dependence.

This behaviour is consistent with the effective sample size diagnostics computed for all monitored parameters: ESS values range from 101 to 2040, with a mean of about 401. Regression coefficients exhibit high sampling efficiency, with effective sample sizes typically ranging between approximately 1500 and 1600 for most parameters. The variance components also show very good mixing, with ESS values of about 1600 for σ^2 and around 1200 for τ , whereas lower ESS values are concentrated among a subset of station-specific effects, which is expected in hierarchical models and does not indicate pathological behavior.

The Geweke diagnostic provides a more conservative assessment. Regression coefficients and vari-

ance components have z -scores within the conventional interval $[-2, 2]$. For the station-specific effects, 27 out of 116 parameters fall outside this range. The Geweke z -scores for this block range from approximately 0.50 to 2.71, with exceedances only on the upper side of the conventional threshold. These exceedances are concentrated among local random effects rather than structural parameters, and therefore suggest lower Monte Carlo stability for some station-level effects rather than a global failure of convergence.

The Heidelberger–Welch diagnostic was also used to assess stationarity and the adequacy of the burn-in period. The stationarity test is passed by all monitored parameters. The largest suggested starting point is iteration 1001, leading to a burn-in of 1000 iterations, which is adopted before thinning the chain. The half-width test is passed by 116 out of 136 parameters; again, the remaining failures are mainly concentrated among station-specific effects and indicate lower precision in estimating some local posterior means, consistently with the ESS and Geweke results.

Table 2: Summary of Geweke and Heidelberger–Welch diagnostics by parameter block.

Parameter block	Geweke min	Geweke max	HW stationarity	HW half-width
β	-1.397	0.290	18/18	17/18
α	0.503	2.709	116/116	97/116
σ^2	0.573	0.573	1/1	1/1
τ	1.412	1.412	1/1	1/1

Overall, the diagnostics indicate satisfactory convergence for the regression coefficients and variance components. Some station-specific effects display lower sampling efficiency and less stable Geweke statistics, but this behaviour is localized and consistent with the hierarchical structure of the model. Since the Heidelberger–Welch stationarity test is passed by all parameters and no persistent autocorrelation is observed, the posterior draws are considered adequate for the subsequent inferential analysis.

4.2 Posterior inference

The posterior estimates shown in Table 3 provide a clear characterization of the main temporal patterns in PM2.5 concentrations.

The reported credible intervals are 95% highest posterior density (HPD) intervals. Effects whose intervals exclude zero are interpreted as credibly different from the baseline category.

The intercept represents the expected value of the log-transformed outcome for the reference category, namely January and Monday, at an average station. Its posterior mean is approximately $\text{r_round}(\text{beta_mean}[1], 3)$, corresponding to a baseline PM2.5 level of about 13.6 on the original scale.

The posterior estimate of the residual variance is approximately 0.324, corresponding to a residual standard deviation of about 0.57. The posterior mean of the group-level scale parameter is approximately 0.298. Hence, the residual standard deviation is larger than the group-level standard deviation, indicating that within-station variability is more pronounced than persistent cross-station heterogeneity.

At the same time, the group-level variation remains non-negligible. This implies that, even after controlling for temporal effects, monitoring stations still differ in their structural pollution profiles.

Table 3: Posterior means, 95% HPD intervals, and approximate percentage effects relative to the baseline category.

Variable	Mean	95% HPD interval	Effect
Intercept (Jan, Mon)	2.609	[2.514, 2.712]	Baseline (~ 13.6)
February	-0.019	[-0.1, 0.073]	-1.8%
March	-0.364	[-0.448, -0.277]	-30.5%
April	-0.617	[-0.701, -0.531]	-46.0%
May	-0.744	[-0.827, -0.655]	-52.5%
June	-0.122	[-0.211, -0.037]	-11.5%
July	-0.600	[-0.684, -0.515]	-45.1%
August	-0.428	[-0.516, -0.345]	-34.8%
September	-0.802	[-0.884, -0.714]	-55.2%
October	-0.607	[-0.69, -0.519]	-45.5%
November	-0.383	[-0.468, -0.3]	-31.8%
December	-0.070	[-0.149, 0.021]	-6.8%

Such differences are consistent with persistent local factors, such as urban density, traffic exposure, industrial activity, topographical conditions, and local meteorological patterns.

From a modeling perspective, this result supports the use of a hierarchical specification. Although much of the variation is driven by within-station dynamics, the station-level random effects capture a relevant component of systematic heterogeneity that would be overlooked in a fully pooled model.

4.2.1 Temporal effects

The temporal coefficients reported in Table 3 are now interpreted in detail. A clear seasonal pattern emerges: relative to January, most months exhibit negative coefficients, indicating lower PM2.5 concentrations during the rest of the year.

The magnitude of these effects is substantively large. March, April, May, July, August, September, October, and November show clear reductions relative to January, with estimated decreases ranging from approximately 30% to more than 55%. The largest reduction is observed in September, with a posterior mean effect of about -55.2%. February and December are statistically indistinguishable from January, as their HPD intervals include zero. June also shows a smaller but credibly negative effect, corresponding to an estimated reduction of about -11.5%. This pattern is consistent with well-known environmental mechanisms, including increased heating emissions and less favorable atmospheric dispersion conditions during colder months.

For compactness, Table 3 reports only the intercept and monthly effects. The omitted day-of-week coefficients are generally small. Most of their 95% HPD intervals include zero, with the exception of Saturday, which shows a small positive effect relative to Monday, approximately equal to 2.6%. This suggests at most a weak day-of-week pattern once seasonal and station-level heterogeneity are accounted for.

4.2.2 Spatial effects

Figure 5 displays the geographical distribution of monitoring stations, with colors representing the posterior mean of the station-specific random effects.

The analysis is exploratory and relies on latitude-based macro-area groups rather than on a formal spatial model. Specifically, macro-areas are defined using latitude thresholds at 44° and 41.5° . Positive values indicate structurally higher PM2.5 levels after controlling for temporal effects.




Figure 5: Spatial distribution of posterior mean station effects across Italy.

The spatial distribution reveals a clear North–South gradient in the station-level effects. Stations located in northern Italy tend to exhibit positive random effects, whereas southern stations display predominantly negative values. Central regions lie between these two patterns, with an average effect close to zero.

Table 4: Posterior summaries of average station effects by macro-area.

Area	Mean	95% HPD interval	Posterior probability
North	0.191	[0.141, 0.243]	1.0000000
Center	-0.019	[-0.073, 0.029]	0.2244444
South	-0.117	[-0.167, -0.065]	0.0000000
North - South	0.308	[0.289, 0.326]	1.0000000
North - Center	0.210	[0.194, 0.227]	1.0000000

This pattern is confirmed by a posterior comparison of average station effects across macro-areas. The posterior mean station effect is positive in the North, approximately 0.19, with a 95% credible interval entirely above zero. By contrast, the South shows a negative average effect, approximately -0.12, with a 95% credible interval entirely below zero. The Center has an average effect close to

zero, and its credible interval includes zero.

The North–Center contrast is also credibly positive, with an estimated difference of about 0.21 on the log scale. Overall, these results support the presence of a marked geographical gradient in station-level baseline pollution.

This evidence is consistent with well-documented geographical differences in pollution dynamics, including higher industrial intensity, traffic exposure, and less favorable atmospheric dispersion conditions in northern regions. At the same time, the dispersion of the station effects highlights substantial within-region heterogeneity, suggesting that local factors remain important determinants of air quality.

Overall, these results highlight the coexistence of broad geographical gradients and substantial local heterogeneity in $\text{PM}_{2.5}$ dynamics.

4.3 Posterior predictive checks

Posterior predictive checks were performed to assess model adequacy. Replicated datasets were generated from the posterior distribution and compared to the observed data.




Figure 6: Posterior predictive check for the log-transformed $\text{PM}_{2.5}$ outcome.




Figure 7: Posterior predictive checks based on summary statistics.

The distribution of replicated means and standard deviations closely matches the observed values (see Figures 6 and 7), with the observed statistics lying well within the distribution of replicated values. This suggests that the model provides an adequate representation of the first two moments of the data.

Furthermore, the posterior predictive density aligns reasonably well with the empirical distribution of log-transformed PM2.5, although minor discrepancies are visible in the peak and tail regions. These differences are consistent with the Gaussian assumption and do not indicate major model misspecification.

Robustness checks based on alternative prior calibrations, reported in Section 3.5.1, confirm that the main posterior conclusions are stable.

5 Conclusion

This paper develops a Bayesian hierarchical model for daily PM2.5 concentrations, explicitly accounting for heterogeneity across monitoring stations and the unbalanced structure of the panel.

The Metropolis-within-Gibbs sampler exhibits satisfactory convergence behavior, as supported by multiple diagnostic tools. Posterior inference appears stable with respect to prior specification, and posterior predictive checks indicate that the model is able to capture the main features of the observed data reasonably well.

From a substantive perspective, the results highlight a pronounced seasonal pattern, with substantially higher PM2.5 levels during winter months. This finding reinforces the relevance of seasonally targeted environmental policies, particularly those addressing emissions related to heating systems and increased urban activity. By contrast, there is little evidence of a strong weekly cycle: most day-of-week effects are small and not credibly different from zero, although Saturday shows a modest positive deviation. This suggests that short-term policy interventions (e.g., day-specific restrictions) may be insufficient unless embedded within broader, structural measures.

In addition, the presence of persistent differences across monitoring stations, together with the identified North–South gradient, points to the importance of geographically differentiated policy approaches. These results are consistent with the idea that local environmental conditions, urban

structure, and regional economic activity jointly shape air pollution dynamics.

Despite these strengths, the analysis is subject to several limitations. First, the study focuses on a single year (2025), which prevents the investigation of longer-term trends and inter-annual variability. Extending the analysis to multiple years would allow for a more robust assessment of persistence, structural changes, and the potential impact of policy interventions over time.

Second, the spatial coverage of the dataset is limited to 116 monitoring stations and is not fully representative of the entire Italian territory. Some regions are underrepresented, which may introduce sampling bias and limit the external validity of the results, particularly with respect to the estimated spatial gradient.

Third, while the hierarchical structure partially mitigates data imbalance, the irregular temporal coverage across stations may still affect inference, especially for locations with substantial missing periods.

Finally, the current specification does not explicitly model spatial dependence across stations, nor does it incorporate meteorological covariates (e.g., temperature, wind, humidity), which are known to play a crucial role in pollution dynamics.

These limitations suggest several directions for future research. Extending the dataset across multiple years and improving spatial coverage would enhance the robustness of the empirical findings. Incorporating spatial correlation structures, such as Gaussian processes or conditional autoregressive (CAR) models, would allow for a more realistic representation of geographic dependence. Additionally, enriching the model with meteorological and environmental covariates could improve both explanatory power and predictive performance.

More flexible likelihood specifications (e.g., heavy-tailed or skewed distributions) may also be considered to better capture extreme pollution episodes. Finally, the framework could be extended to a fully spatio-temporal hierarchical model, providing a unified approach to jointly model temporal dynamics and spatial interactions.

Overall, the hierarchical Bayesian framework provides a flexible and effective approach for the analysis of air pollution data, offering both methodological rigor and policy-relevant insights.